

Creating Numbers/images with AI: A Hands-on Diffusion Model Exercise

Introduction

In this assignment, you'll learn how to create an AI model that can generate realistic images from scratch using a powerful technique called 'diffusion'. Think of it like teaching AI to draw by first learning how images get blurry and then learning to make them clear again.

What We'll Build

- A diffusion model capable of generating realistic images
- For most students: An AI that generates handwritten digits (0-9) using the MNIST **dataset**
- For students with more computational resources: Options to work with more complex **datasets**
- Visual demonstrations of how random noise gradually transforms into clear, recognizable images
- By the end, your AI should create images realistic enough for another AI to recognize them

Dataset Options

This lab offers flexibility based on your available computational resources:

- Standard Option (Free Colab): We'll primarily use the MNIST handwritten digit **dataset**, which works well with limited GPU memory and completes training in a reasonable time frame. Most examples and code in this notebook are optimized for MNIST.
- Advanced Option: If you have access to more powerful GPUs (either through Colab Pro/Pro+ or your own hardware), you can experiment with more complex **datasets** like Fashion-MNIST, CIFAR-10, or even face generation. You'll need to adapt the model architecture, hyperparameters, and evaluation metrics accordingly.

Resource Requirements

- Basic MNIST: Works with free Colab GPUs (2-4GB VRAM), ~30 minutes training
- Fashion-MNIST: Similar requirements to MNIST CIFAR-10: Requires more memory (8-12GB VRAM) and longer training (~2 hours)
- Higher resolution images: Requires substantial GPU resources and several hours of training

Before You Start

1. Make sure you're running this in Google Colab or another environment with GPU access
2. Go to 'Runtime' → 'Change runtime type' and select 'GPU' as your hardware accelerator
3. Each code cell has comments explaining what it does
4. Don't worry if you don't understand every detail - focus on the big picture!
5. If working with larger **datasets**, monitor your GPU memory usage carefully

The concepts you learn with MNIST will scale to more complex **datasets**, so even if you're using the basic option, you'll gain valuable knowledge about generative AI that applies to more advanced applications.

Step 1: Setting Up Our Tools

First, let's install and import all the tools we need. Run this cell and wait for it to complete.

```
1 # Step 1: Install required packages
2 %pip install einops
3 print("Package installation complete.")
4
5 # Step 2: Import libraries
6 # --- Core PyTorch libraries ---
7 import torch # Main deep learning framework
8 import torch.nn.functional as F # Neural network functions like activation functions
9 import torch.nn as nn # Neural network building blocks (layers)
10 from torch.optim import Adam # Optimization algorithm for training
11
12 # --- Data handling ---
13 from torch.utils.data import Dataset, DataLoader # For organizing and loading our data
14 import torchvision # Library for computer vision datasets and models
15 import torchvision.transforms as transforms # For preprocessing images
16
17 # --- Tensor manipulation ---
18 import random # For random operations
19 from einops.layers.torch import Rearrange # For reshaping tensors in neural networks
```

```

20 from einops import rearrange # For elegant tensor reshaping operations
21 import numpy as np # For numerical operations on arrays
22
23 # --- System utilities ---
24 import os # For operating system interactions (used for CPU count)
25
26 # --- Visualization tools ---
27 import matplotlib.pyplot as plt # For plotting images and graphs
28 from PIL import Image # For image processing
29 from torchvision.utils import save_image, make_grid # For saving and displaying image grids
30
31 # Step 3: Set up device (GPU or CPU)
32 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
33 print(f"We'll be using: {device}")
34
35 # Check if we're actually using GPU (for students to verify)
36 if device.type == "cuda":
37     print(f"GPU name: {torch.cuda.get_device_name(0)}")
38     print(f"GPU memory: {torch.cuda.get_device_properties(0).total_memory / 1e9:.2f} GB")
39 else:
40     print("Note: Training will be much slower on CPU. Consider using Google Colab with GPU enabled.")

```

Requirement already satisfied: einops in /usr/local/lib/python3.12/dist-packages (0.8.1)
 Package installation complete.
 We'll be using: cuda
 GPU name: Tesla T4
 GPU memory: 15.83 GB

REPRODUCIBILITY AND DEVICE SETUP

```

1 # Step 4: Set random seeds for reproducibility
2 # Diffusion models are sensitive to initialization, so reproducible results help with debugging
3 SEED = 42 # Universal seed value for reproducibility
4 torch.manual_seed(SEED)           # PyTorch random number generator
5 np.random.seed(SEED)             # NumPy random number generator
6 random.seed(SEED)               # Python's built-in random number generator
7
8 print(f"Random seeds set to {SEED} for reproducible results")
9
10 # Configure CUDA for GPU operations if available
11 if torch.cuda.is_available():
12     torch.cuda.manual_seed(SEED)      # GPU random number generator
13     torch.cuda.manual_seed_all(SEED) # All GPUs random number generator
14
15     # Ensure deterministic GPU operations
16     # Note: This slightly reduces performance but ensures results are reproducible
17     torch.backends.cudnn.deterministic = True
18     torch.backends.cudnn.benchmark = False
19
20     try:
21         # Check available GPU memory
22         gpu_memory = torch.cuda.get_device_properties(0).total_memory / 1e9 # Convert to GB
23         print(f"Available GPU Memory: {gpu_memory:.1f} GB")
24
25         # Add recommendation based on memory
26         if gpu_memory < 4:
27             print("Warning: Low GPU memory. Consider reducing batch size if you encounter OOM errors.")
28     except Exception as e:
29         print(f"Could not check GPU memory: {e}")
30 else:
31     print("No GPU detected. Training will be much slower on CPU.")
32     print("If you're using Colab, go to Runtime > Change runtime type and select GPU.")

```

Random seeds set to 42 for reproducible results
 Available GPU Memory: 15.8 GB

Step 2: Choosing Your Dataset

You have several options for this exercise, depending on your computer's capabilities:

Option 1: MNIST (Basic - Works on Free Colab)

- Content: Handwritten digits (0-9)

- Image size: 28x28 pixels, Grayscale
- Training samples: 60,000
- Memory needed: ~2GB GPU
- Training time: ~15-30 minutes on Colab
- **Choose this if:** You're using free Colab or have a basic GPU

Option 2: Fashion-MNIST (Intermediate)

- Content: Clothing items (shirts, shoes, etc.)
- Image size: 28x28 pixels, Grayscale
- Training samples: 60,000
- Memory needed: ~2GB GPU
- Training time: ~15-30 minutes on Colab
- **Choose this if:** You want more interesting images but have limited GPU

Option 3: CIFAR-10 (Advanced)

- Content: Real-world objects (cars, animals, etc.)
- Image size: 32x32 pixels, Color (RGB)
- Training samples: 50,000
- Memory needed: ~4GB GPU
- Training time: ~1-2 hours on Colab
- **Choose this if:** You have Colab Pro or a good local GPU (8GB+ memory)

Option 4: CelebA (Expert)

- Content: Celebrity face images
- Image size: 64x64 pixels, Color (RGB)
- Training samples: 200,000
- Memory needed: ~8GB GPU
- Training time: ~3-4 hours on Colab
- **Choose this if:** You have excellent GPU (12GB+ memory)

To use your chosen **dataset**, uncomment its section in the code below and make sure all others are commented out.

```

1 #=====
2 # SECTION 2: DATASET SELECTION AND CONFIGURATION
3 #=====
4 # STUDENT INSTRUCTIONS:
5 # 1. Choose ONE dataset option based on your available GPU memory
6 # 2. Uncomment ONLY ONE dataset section below
7 # 3. Make sure all other dataset sections remain commented out
8
9 #-----
10 # OPTION 1: MNIST (Basic - 2GB GPU)
11 #-----
12 # Recommended for: Free Colab or basic GPU
13 # Memory needed: ~2GB GPU
14 # Training time: ~15-30 minutes
15
16 IMG_SIZE = 28
17 IMG_CH = 1
18 N_CLASSES = 10
19 BATCH_SIZE = 64
20 EPOCHS = 30
21
22 import torchvision # Added import
23 import torchvision.transforms as transforms
24 from torch.utils.data import DataLoader # Added import
25 import os # Added import
26
27 transform = transforms.Compose([
28     transforms.ToTensor(),
29     transforms.Normalize((0.5,), (0.5,))
30 ])
31
32 # Your code to load the MNIST dataset
33 # Hint: Use torchvision.datasets.MNIST with root='./data', train=True,
34 #        transform=transform, and download=True
35 # Then print a success message
36

```

```

36
37 # Enter your code here:
38 # --- Load MNIST dataset ---
39 train_dataset = torchvision.datasets.MNIST(
40     root='./data',
41     train=True,
42     transform=transform,
43     download=True
44 )
45
46 test_dataset = torchvision.datasets.MNIST(
47     root='./data',
48     train=False,
49     transform=transform,
50     download=True
51 )
52
53 # --- Create DataLoaders ---
54 train_loader = DataLoader(train_dataset, batch_size=BATCH_SIZE, shuffle=True, num_workers=os.cpu_count())
55 test_loader = DataLoader(test_dataset, batch_size=BATCH_SIZE, shuffle=False, num_workers=os.cpu_count())
56
57 print(f"✅ MNIST dataset loaded successfully! Total training samples: {len(train_dataset)}")
58 print(f"Example image shape: {train_dataset[0][0].shape}")

```

✅ MNIST dataset loaded successfully! Total training samples: 60000
Example image shape: torch.Size([1, 28, 28])

```

1 =====
2 # VALIDATING DATASET SELECTION & GPU MEMORY (MNIST ONLY)
3 =====
4
5 # --- Helper: probe presence of common dataset vars ---
6 has_explicit_flag = 'dataset' in globals()
7 has_train_dataset = 'train_dataset' in globals()
8 has_test_dataset = 'test_dataset' in globals()
9 has_train_loader = 'train_loader' in globals()
10 has_test_loader = 'test_loader' in globals()
11
12 if not (has_explicit_flag or has_train_dataset or has_test_dataset or has_train_loader or has_test_loader):
13     raise ValueError("")
14 ✖ ERROR: No dataset selected! Please uncomment the MNIST dataset option.
15 Available option:
16 1. MNIST (Basic) - 2GB GPU
17 """
18
19 # --- Infer dataset name if not explicitly set ---
20 def _infer_dataset_name():
21     if 'dataset' in globals() and isinstance(dataset, str):
22         return dataset
23     if 'train_dataset' in globals():
24         return train_dataset.__class__.__name__
25     if 'train_loader' in globals():
26         try:
27             return train_loader.dataset.__class__.__name__
28         except Exception:
29             pass
30     if 'IMG_SIZE' in globals() and 'IMG_CH' in globals():
31         if IMG_SIZE == 28 and IMG_CH == 1:
32             return "MNIST"
33     return "Unknown"
34
35 dataset_name = _infer_dataset_name()
36
37 if dataset_name == "MNIST" and 'dataset' not in globals():
38     dataset = "MNIST"
39
40 print(f"📦 Detected dataset: {dataset_name}")
41
42 # --- GPU memory requirement (MNIST only) ---
43 required_gb = 2
44
45 if torch.cuda.is_available():
46     props = torch.cuda.get_device_properties(0)
47     gpu_gb = props.total_memory / 1e9
48     print(f"💻 GPU: {props.name}")
49     print(f"🧠 VRAM: {gpu_gb:.2f} GB (recommended ≥ {required_gb} GB for MNIST)")
50     if gpu_gb < required_gb:

```

```

51     print("⚠️ VRAM may be insufficient - consider reducing batch size.")
52 else:
53     print("✅ GPU memory check passed for MNIST.")
54 else:
55     print("⚠️ No GPU detected. Training will be slow on CPU. Enable GPU in Colab if possible.")
56

```

Detected dataset: MNIST
 GPU: Tesla T4
 VRAM: 15.83 GB (recommended ≥ 2 GB for MNIST)
 GPU memory check passed for MNIST.

```

1 # --- Sample batch properties ---
2 _probe_loader = DataLoader(train_dataset, batch_size=16, shuffle=True, num_workers=0)
3 xb, yb = next(iter(_probe_loader)) # xb: [B, C, H, W], yb: [B]
4
5 print("Batch shape:", tuple(xb.shape))
6 print("DType:", xb.dtype, "| Labels dtype:", yb.dtype)
7 print("Image size (C,H,W):", tuple(xb.shape[1:]))
8 print("Value range (min,max):", float(xb.min()), float(xb.max()))
9 print("Mean/std (first batch):", xb.mean().item(), xb.std().item())
10 print("Unique labels in batch:", torch.unique(yb).tolist())
11
12 # Optional: quick visual sanity check (comment out if running headless)
13 # grid = make_grid(xb[:16], nrow=8, normalize=True, value_range=(-1, 1), padding=1)
14 # plt.figure(figsize=(6,3))
15 # plt.imshow(grid.permute(1, 2, 0).cpu().numpy(), cmap="gray")
16 # plt.axis("off")
17 # plt.title("MNIST sample batch")
18 # plt.show()
19

```

```

Batch shape: (16, 1, 28, 28)
DType: torch.float32 | Labels dtype: torch.int64
Image size (C,H,W): (1, 28, 28)
Value range (min,max): -1.0 1.0
Mean/std (first batch): -0.7353366017341614 0.6148757338523865
Unique labels in batch: [0, 1, 2, 3, 4, 5, 6, 8, 9]

```

▼ Step 3: Building Our Model Components

Now we'll create the building blocks of our AI model. Think of these like LEGO pieces that we'll put together to make our number generator:

- GELUConvBlock: The basic building block that processes images
- DownBlock: Makes images smaller while finding important features
- UpBlock: Makes images bigger again while keeping the important features
- Other blocks: Help the model understand time and what number to generate

```

1 # Basic building block that processes images
2 class GELUConvBlock(nn.Module):
3     def __init__(self, in_ch, out_ch, group_size):
4         """
5             Creates a block with convolution, normalization, and activation
6
7             Args:
8                 in_ch (int): Number of input channels
9                 out_ch (int): Number of output channels
10                group_size (int): Number of groups for GroupNorm
11            """
12     super().__init__()
13
14     # Check that group_size is compatible with out_ch
15     if out_ch % group_size != 0:
16         print(f"Warning: out_ch ({out_ch}) is not divisible by group_size ({group_size})")
17     # Adjust group_size to be compatible
18     group_size = min(group_size, out_ch)
19     while out_ch % group_size != 0:
20         group_size -= 1
21     print(f"Adjusted group_size to {group_size}")
22
23     # --- Layers: Conv → GroupNorm → GELU ---
24     self.block = nn.Sequential(

```

```

25         nn.Conv2d(in_ch, out_ch, kernel_size=3, padding=1), # preserves spatial size
26         nn.GroupNorm(num_groups=group_size, num_channels=out_ch),
27         nn.GELU()
28     )
29
30     def forward(self, x):
31         # Pass input through the sequential block
32         return self.block(x)
33

```



```

1 # Rearranges pixels to downsample the image (2x reduction in spatial dimensions)
2 class RearrangePoolBlock(nn.Module):
3     def __init__(self, in_chs, group_size):
4         """
5             Downsamples the spatial dimensions by 2x while preserving information
6
7             Args:
8                 in_chs (int): Number of input channels
9                 group_size (int): Number of groups for GroupNorm
10            """
11            super().__init__()
12            # Space-to-depth: (B, C, H, W) -> (B, 4C, H/2, W/2)
13            self.rearrange = Rearrange('b c (h s1) (w s2) -> b (c s1 s2) h w', s1=2, s2=2)
14
15            # After rearrange, channels become 4 * in_chs.
16            # Typical U-Net encoder doubles channels at each downsample:
17            out_chs = in_chs * 2
18            # Fix: The input channels to GELUConvBlock should match the output channels of rearrange (in_chs * 4)
19            self.conv = GELUConvBlock(in_chs * 4, out_chs, group_size)
20
21     def forward(self, x):
22         x = self.rearrange(x)    # downsample H,W by 2; multiply channels by 4
23         x = self.conv(x)        # process features
24         return x

```

```

1 #Let's implement the upsampling block for our U-Net architecture:
2 import torch.nn as nn # Import nn module
3 import torch.nn.functional as F # Import F module
4
5 class DownBlock(nn.Module):
6     """
7         Downsampling block for encoding path in U-Net architecture.
8
9         This block:
10            1. Processes input features with two convolutional blocks to change channels
11            2. Downsamples spatial dimensions by 2x using a strided convolution
12
13         Args:
14             in_chs (int): Number of input channels
15             out_chs (int): Number of output channels
16             group_size (int): Number of groups for GroupNorm
17     """
18     def __init__(self, in_chs, out_chs, group_size):
19         super().__init__()
20
21         # --- Layers: Conv → GN → GELU → Conv → GN → GELU → Downsample ---
22         self.model = nn.Sequential(
23             GELUConvBlock(in_chs, out_chs, group_size), # First conv block changes channel dimensions
24             GELUConvBlock(out_chs, out_chs, group_size), # Second conv block processes features
25             # Spatial downsampling using a strided convolution
26             nn.Conv2d(out_chs, out_chs, kernel_size=3, stride=2, padding=1)
27         )
28
29         # Log the configuration for debugging
30         print(f"Created DownBlock: in_chs={in_chs}, out_chs={out_chs}, spatial_reduction=2x")
31
32     def forward(self, x):
33         """
34             Forward pass through the DownBlock.
35
36             Args:
37                 x (torch.Tensor): Input tensor of shape [B, in_chs, H, W]
38
39             Returns:
40                 torch.Tensor: Output tensor of shape [B, out_chs, H/2, W/2]

```

```

1 # Now let's implement the upsampling block for our U-Net architecture:
2 class UpBlock(nn.Module):
3     """
4         Upsampling block for decoding path in U-Net architecture.
5
6         This block:
7             1. Takes features from the decoding path and corresponding skip connection
8             2. Upsamples spatial dimensions by 2x using transposed convolution
9             3. Concatenates them along the channel dimension
10            4. Processes features through multiple convolutional blocks
11
12     Args:
13         in_chs (int): Number of input channels from the previous layer (decoder)
14         out_chs (int): Number of output channels for this stage
15         group_size (int): Number of groups for GroupNorm
16     """
17     def __init__(self, in_chs, out_chs, group_size):
18         super().__init__()
19
20         # 2x upsampling via transposed convolution: keeps channels at in_chs
21         self.up = nn.ConvTranspose2d(in_chs, in_chs, kernel_size=2, stride=2)
22
23         # After upsample, we concatenate with skip (assume skip has in_chs channels in a symmetric U-Net)
24         fused_chs = in_chs * 2
25
26         # Two Conv→GN→GELU blocks to refine fused features
27         self.refine = nn.Sequential(
28             GELUConvBlock(fused_chs, out_chs, group_size),
29             GELUConvBlock(out_chs, out_chs, group_size)
30         )
31
32         # Log the configuration for debugging
33         print(f"Created UpBlock: in_chs={in_chs}, out_chs={out_chs}, spatial_increase=2x")
34
35     def forward(self, x, skip):
36         """
37             Forward pass through the UpBlock.
38
39         Args:
40             x (torch.Tensor): Decoder input [B, in_chs, H, W]
41             skip (torch.Tensor): Encoder skip [B, in_chs, 2H, 2W]
42
43         Returns:
44             torch.Tensor: [B, out_chs, 2H, 2W]
45         """
46
47         # 1) Upsample decoder features to match skip spatial size
48         x = self.up(x) # -> [B, in_chs, 2H, 2W]
49
50         # Handle potential off-by-one mismatches (odd sizes) with pad/crop
51         if x.shape[-2:] != skip.shape[-2:]:
52             dh = skip.shape[-2] - x.shape[-2]
53             dw = skip.shape[-1] - x.shape[-1]
54             # pad (right/bottom) if smaller
55             x = F.pad(x, (0, max(0, dw), 0, max(0, dh)))
56             # crop if we overshot
57             x = x[..., :skip.shape[-2], :skip.shape[-1]]
58
59         # 2) Concatenate along channels
60         x = torch.cat([x, skip], dim=1) # [B, in_chs*2, 2H, 2W]
61
62         # 3) Refine fused features
63         return self.refine(x)

```

```
1 # Here we implement the time embedding block for our U-Net architecture:  
2 # Helps the model understand time steps in diffusion process  
3 class SinusoidalPositionEmbedBlock(nn.Module):  
4     """  
5         Creates sinusoidal embeddings for time steps in diffusion process.  
6  
7         This embedding scheme is adapted from the Transformer architecture and  
8         provides a unique representation for each time step that preserves  
9         relative distance information.  
10    """
```

```

10
11     Args:
12         dim (int): Embedding dimension
13     """
14     def __init__(self, dim):
15         super().__init__()
16         self.dim = dim
17
18     def forward(self, time):
19         """
20             Computes sinusoidal embeddings for given time steps.
21
22         Args:
23             time (torch.Tensor): Time steps tensor of shape [batch_size]
24
25         Returns:
26             torch.Tensor: Time embeddings of shape [batch_size, dim]
27         """
28         device = time.device
29         half_dim = self.dim // 2
30         embeddings = torch.log(torch.tensor(10000.0, device=device)) / (half_dim - 1)
31         embeddings = torch.exp(torch.arange(half_dim, device=device) * -embeddings)
32         embeddings = time[:, None] * embeddings[None, :]
33         embeddings = torch.cat((embeddings.sin(), embeddings.cos()), dim=-1)
34
35     return embeddings
36

```

```

1 class EmbedBlock(nn.Module):
2     """
3         Creates embeddings for class conditioning in diffusion models.
4
5         This module transforms a one-hot or index representation of a class
6         into a rich embedding that can be added to feature maps.
7
8     Args:
9         input_dim (int): Input dimension (typically number of classes)
10        emb_dim (int): Output embedding dimension
11    """
12    def __init__(self, input_dim, emb_dim):
13        super(EmbedBlock, self).__init__()
14        self.input_dim = input_dim
15
16        # --- Class embedding model ---
17        self.model = nn.Sequential(
18            nn.Linear(input_dim, emb_dim),
19            nn.GELU(),
20            nn.Linear(emb_dim, emb_dim),
21            nn.GELU(),
22            nn.Unflatten(1, (emb_dim, 1, 1)) # reshape for broadcasting with feature maps
23        )
24
25    def forward(self, x):
26        """
27            Computes class embeddings for the given class indices.
28
29        Args:
30            x (torch.Tensor): Class indices [batch_size] or one-hot encodings [batch_size, input_dim]
31
32        Returns:
33            torch.Tensor: Class embeddings [batch_size, emb_dim, 1, 1]
34        """
35        # If input is class indices, convert to one-hot
36        if x.ndim == 1 or x.shape[1] == 1:
37            x = F.one_hot(x.long().squeeze(), num_classes=self.input_dim).float()
38
39        x = x.view(-1, self.input_dim)
40
41        return self.model(x)

```

```

1 # Main U-Net model that puts everything together
2 class UNet(nn.Module):
3     """
4         U-Net architecture for diffusion models with time and class conditioning.
5     """
6     def __init__(self, T, img_ch, img_size, down_chs, t_embed_dim, c_embed_dim):

```

```

7     super().__init__()
8     self.T = T
9     self.img_ch = img_ch
10    self.img_size = img_size
11    self.down_chs = down_chs
12    self.group_size = 8 # default GN group size (GELUConvBlock will auto-fix if not divisible)
13
14    base_ch = down_chs[0]
15    deepest_ch = down_chs[-1]
16
17    # ---- Time embedding: sinusoidal -> MLP -> [B, deepest_ch, 1, 1]
18    self.time_mlp = nn.Sequential(
19        SinusoidalPositionEmbedBlock(t_embed_dim),
20        nn.Linear(t_embed_dim, t_embed_dim * 4),
21        nn.GELU(),
22        nn.Linear(t_embed_dim * 4, deepest_ch),
23        nn.Unflatten(1, (deepest_ch, 1, 1)),
24    )
25
26    # ---- Class embedding: one-hot / index -> MLP -> project to deepest_ch -> [B, deepest_ch, 1, 1]
27    # Uses global N_CLASSES defined earlier
28    self.class_embed = EmbedBlock(input_dim=N_CLASSES, emb_dim=c_embed_dim)
29    self.class_to_ch = nn.Conv2d(c_embed_dim, deepest_ch, kernel_size=1)
30
31    # ---- Initial feature extraction
32    self.in_conv = GELUConvBlock(img_ch, base_ch, self.group_size)
33
34    # ---- Down path: list of DownBlocks
35    downs = []
36    for i in range(len(down_chs) - 1):
37        downs.append(DownBlock(down_chs[i], down_chs[i + 1], self.group_size))
38    self.down_blocks = nn.ModuleList(downs)
39
40    # ---- Mid (bottleneck) blocks
41    self.mid1 = GELUConvBlock(deepest_ch, deepest_ch, self.group_size)
42    self.mid2 = GELUConvBlock(deepest_ch, deepest_ch, self.group_size)
43
44    # ---- Up path: list of UpBlocks (reverse order)
45    ups = []
46    for i in reversed(range(len(down_chs) - 1)):
47        # Decoder input channels at this stage = channels from deeper level
48        dec_in = down_chs[i + 1]
49        dec_out = down_chs[i]
50        ups.append(UpBlock(in_chs=dec_in, out_chs=dec_out, group_size=self.group_size))
51    self.up_blocks = nn.ModuleList(ups)
52
53    # To make concatenation work even if skip channels != decoder in_chs,
54    # create 1x1 projections for each skip (align to decoder in_chs at that stage).
55    # Order matches self.up_blocks (deepest -> shallow).
56    skip_projs = []
57    for i in reversed(range(len(down_chs) - 1)):
58        skip_chs = down_chs[i] # channels before DownBlock at level i
59        target_chs = down_chs[i + 1] # decoder in_chs used by UpBlock at that level
60        skip_projs.append(nn.Conv2d(skip_chs, target_chs, kernel_size=1))
61    self.skip_projs = nn.ModuleList(skip_projs)
62
63    # ---- Final projection back to image channels
64    self.out_conv = nn.Conv2d(base_ch, img_ch, kernel_size=1)
65
66    print(f"Created UNet with {len(down_chs)} scale levels")
67    print(f"Channel dimensions: {down_chs}")
68
69    def forward(self, x, t, c, c_mask):
70        """
71            x:      [B, img_ch, H, W]           (noisy image)
72            t:      [B]                      (timestep indices)
73            c:      [B] or [B, N_CLASSES]    (class labels or one-hot)
74            c_mask: [B, 1]                  (1=use class, 0=ignore; for classifier-free guidance)
75        """
76        B = x.size(0)
77
78        # ---- Time embedding
79        t_emb = self.time_mlp(t) # [B, deepest_ch, 1, 1]
80
81        # ---- Class embedding (supports indices or one-hot)
82        c_emb = self.class_embed(c) # [B, c_embed_dim, 1, 1]
83        c_emb = self.class_to_ch(c_emb) # [B, deepest_ch, 1, 1]

```

```

84     # Apply mask (broadcast)
85     if c_mask is not None:
86         c_mask_ = c_mask.view(B, 1, 1, 1).to(c_emb.device).type_as(c_emb)
87         c_emb = c_emb * c_mask_
88
89     # ---- Initial conv
90     h = self.in_conv(x)                      # [B, down_chs[0], H, W]
91
92     # ---- Down path with skip collection
93     skips = []
94     for i, down in enumerate(self.down_blocks):
95         # Save skip BEFORE downsample (so its spatial size is larger; matches UpBlock's expectation)
96         skips.append(h)                      # [B, down_chs[i], H_i, W_i]
97         h = down(h)                        # [B, down_chs[i+1] (=), H_i/2, W_i/2]
98
99     # ---- Mid (inject conditioning)
100    h = self.mid1(h)
101    h = h + t_emb + c_emb                # broadcast add at bottleneck
102    h = self.mid2(h)
103
104    # ---- Up path (reverse order), align skip channels, then fuse
105    for idx, up in enumerate(self.up_blocks):
106        skip = skips[-(idx + 1)]          # matching skip (higher-res)
107        # project skip channels to match decoder in_chs at this level
108        skip_aligned = self.skip_projs[idx](skip)
109        h = up(h, skip_aligned)          # [B, out_chs(level), 2H, 2W]
110
111    # ---- Final projection to image channels (predict noise)
112    out = self.out_conv(h)              # [B, img_ch, H, W]
113
114    return out

```

▼ Step 4: Setting Up The Diffusion Process

Now we'll create the process of adding and removing noise from images. Think of it like:

1. Adding fog: Slowly making the image more and more blurry until you can't see it
2. Removing fog: Teaching the AI to gradually make the image clearer
3. Controlling the process: Making sure we can generate specific numbers we want

```

1 import torch
2
3 # Set up device (GPU or CPU) - Added to ensure 'device' is defined
4 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
5 print(f"Using device: {device}")
6
7 # Set up the noise schedule
8 n_steps = 100 # How many steps to go from clear image to noise
9 beta_start = 0.0001 # Starting noise level (small)
10 beta_end = 0.02      # Ending noise level (larger)
11
12 # Create schedule of gradually increasing noise levels
13 beta = torch.linspace(beta_start, beta_end, n_steps).to(device)
14
15 # Calculate important values used in diffusion equations
16 alpha = 1 - beta # Portion of original image to keep at each step
17 alpha_bar = torch.cumprod(alpha, dim=0) # Cumulative product of alphas
18 sqrt_alpha_bar = torch.sqrt(alpha_bar) # For scaling the original image
19 sqrt_one_minus_alpha_bar = torch.sqrt(1 - alpha_bar) # For scaling the noise

```

Using device: cuda

```

1
2 # Function to add noise to images (forward diffusion process)
3 def add_noise(x_0, t):
4     """
5     Add noise to images according to the forward diffusion process.
6
7     The formula is: x_t = \sqrt(\alpha_{bar_t}) * x_0 + \sqrt(1-\alpha_{bar_t}) * \epsilon
8     where \epsilon is random noise and \alpha_{bar_t} is the cumulative product of (1-\beta).
9
10    Args:
11        x_0 (torch.Tensor): Original clean image [B, C, H, W]

```

```

12     t (torch.Tensor): Timestep indices indicating noise level [B]
13
14     Returns:
15         tuple: (noisy_image, noise_added)
16             - noisy_image is the image with noise added
17             - noise_added is the actual noise that was added (for training)
18     """
19     # Create random Gaussian noise with same shape as image
20     noise = torch.randn_like(x_0)
21
22     # Get noise schedule values for the specified timesteps
23     # Reshape to allow broadcasting with image dimensions
24     sqrt_alpha_bar_t = sqrt_alpha_bar[t].reshape(-1, 1, 1, 1)
25     sqrt_one_minus_alpha_bar_t = sqrt_one_minus_alpha_bar[t].reshape(-1, 1, 1, 1)
26
27     # Apply the forward diffusion equation:
28     # Mixture of original image (scaled down) and noise (scaled up)      # Your code to apply the forward diffusion equation
29     # Hint: Mix the original image and noise according to the noise schedule
30
31     # Enter your code here:
32     x_t = sqrt_alpha_bar_t * x_0 + sqrt_one_minus_alpha_bar_t * noise
33
34     return x_t, noise

```

```

1 # Function to remove noise from images (reverse diffusion process)
2 @torch.no_grad()  # Don't track gradients during sampling (inference only)
3 def remove_noise(x_t, t, model, c, c_mask):
4     """
5     Remove noise from images using the learned reverse diffusion process.
6
7     This implements a single step of the reverse diffusion sampling process.
8     The model predicts the noise in the image, which we then use to partially
9     denoise the image.
10
11    Args:
12        x_t (torch.Tensor): Noisy image at timestep t [B, C, H, W]
13        t (torch.Tensor): Current timestep indices [B]
14        model (nn.Module): U-Net model that predicts noise
15        c (torch.Tensor): Class conditioning (what digit to generate) [B, C]
16        c_mask (torch.Tensor): Mask for conditional generation [B, 1]
17
18    Returns:
19        torch.Tensor: Less noisy image for the next timestep [B, C, H, W]
20    """
21    # Predict the noise in the image using our model
22    predicted_noise = model(x_t, t, c, c_mask)
23
24    # Get noise schedule values for the current timestep
25    alpha_t = alpha[t].reshape(-1, 1, 1, 1)
26    alpha_bar_t = alpha_bar[t].reshape(-1, 1, 1, 1)
27    beta_t = beta[t].reshape(-1, 1, 1, 1)
28    # Correctly calculate sqrt_one_minus_alpha_bar_t
29    sqrt_one_minus_alpha_bar_t = sqrt_one_minus_alpha_bar[t].reshape(-1, 1, 1, 1)
30
31    # Special case: if we're at the first timestep (t=0), we're done
32    if t[0] == 0:
33        return x_t
34    else:
35        # Calculate the mean of the denoised distribution
36        # This is derived from Bayes' rule and the diffusion process equations
37        mean = (1 / torch.sqrt(alpha_t)) * (
38            x_t - (beta_t / sqrt_one_minus_alpha_bar_t) * predicted_noise
39        )
40
41        # Add a small amount of random noise (variance depends on timestep)
42        # This helps prevent the generation from becoming too deterministic
43        noise = torch.randn_like(x_t)
44
45        # Return the partially denoised image with a bit of new random noise
46        return mean + torch.sqrt(beta_t) * noise

```

```

1 @torch.no_grad()
2 def show_noise_progression(image, num_steps=5):
3     """
4     Visualize how an image gets progressively noisier in the diffusion process.

```

```

5
6     Args:
7         image (torch.Tensor): [C, H, W] tensor (normalized to [-1, 1] if you used the MNIST transform)
8         num_steps (int): number of panels to show (including the original)
9     """
10    assert image.ndim == 3, "Expected image of shape [C, H, W]"
11    img = image.to(device)
12
13    plt.figure(figsize=(3 * num_steps, 3))
14
15    # Helper to rescale for display
16    def _display_ready(x):
17        x = x.detach().float().cpu()
18        if x.min() < 0: # handle [-1, 1] normalization
19            x = (x + 1) / 2
20        return x.clamp(0, 1)
21
22    # Panel 1: original
23    plt.subplot(1, num_steps, 1)
24    if IMG_CH == 1:
25        plt.imshow(_display_ready(img)[0], cmap='gray')
26    else:
27        plt.imshow(_display_ready(img).permute(1, 2, 0))
28    plt.title('Original')
29    plt.axis('off')
30
31    # Noisier versions
32    for i in range(1, num_steps):
33        # pick an evenly spaced timestep
34        t_idx = int(round((i / (num_steps - 1)) * (n_steps - 1))) if num_steps > 1 else 0
35        t_idx = max(0, min(n_steps - 1, t_idx))
36        t = torch.tensor([t_idx], device=device)
37
38        noisy, _ = add_noise(img.unsqueeze(0), t) # [1, C, H, W]
39        noisy_img = noisy[0]
40
41        plt.subplot(1, num_steps, i + 1)
42        if IMG_CH == 1:
43            plt.imshow(_display_ready(noisy_img)[0], cmap='gray')
44        else:
45            plt.imshow(_display_ready(noisy_img).permute(1, 2, 0))
46        pct = int(round(100 * t_idx / (n_steps - 1))) if n_steps > 1 else 0
47        plt.title(f't={t_idx} ({~{pct}% noise})')
48        plt.axis('off')
49
50    plt.tight_layout()
51    plt.show()
52

```

▼ Step 5: Training Our Model

Now we'll teach our AI to generate images. This process:

1. Takes a clear image
2. Adds random noise to it
3. Asks our AI to predict what noise was added
4. Helps our AI learn from its mistakes

This will take a while, but we'll see progress as it learns!

```

1 # Create our model and move it to GPU if available
2 model = UNet(
3     T=n_steps, # Number of diffusion time steps
4     img_ch=IMG_CH, # Number of channels in our images (1 for grayscale, 3 for RGB)
5     img_size=IMG_SIZE, # Size of input images (28 for MNIST, 32 for CIFAR-10)
6     down_chs=(32, 64, 128), # Channel dimensions for each downsampling level
7     t_embed_dim=8, # Dimension for time step embeddings
8     c_embed_dim=N_CLASSES # Number of classes for conditioning
9 ).to(device)
10
11 # Print model summary
12 print(f"\n{'*'*50}")
13 print(f"MODEL ARCHITECTURE SUMMARY")
14 print(f"{'*'*50}")

```

```

15 print(f"Input resolution: {IMG_SIZE}x{IMG_SIZE}")
16 print(f"Input channels: {IMG_CH}")
17 print(f"Time steps: {n_steps}")
18 print(f"Condition classes: {N_CLASSES}")
19 print(f"GPU acceleration: {'Yes' if device.type == 'cuda' else 'No'}")
20
21 # Validate model parameters and estimate memory requirements
22 # Hint: Create functions to count parameters and estimate memory usage
23
24 # Enter your code here:
25 # ---- Model stats: params and rough VRAM estimate ----
26 def count_parameters(m):
27     total = sum(p.numel() for p in m.parameters())
28     trainable = sum(p.numel() for p in m.parameters() if p.requires_grad)
29     return total, trainable
30
31 def human_bytes(n):
32     for u in ["B", "KB", "MB", "GB", "TB"]:
33         if n < 1024: return f"{n:.2f} {u}"
34         n /= 1024
35     return f"{n:.2f} PB"
36
37 def estimate_param_memory_bytes(m, dtype_bytes=4):
38     # fp32 training rough rule of thumb:
39     # weights + grads + Adam m + Adam v ~ 4 copies
40     n_params = sum(p.numel() for p in m.parameters())
41     return n_params * 4 * dtype_bytes
42
43 total_params, trainable_params = count_parameters(model)
44 param_mem_bytes = estimate_param_memory_bytes(model)
45
46 print(f"\nParameters: {total_params:,} (trainable: {trainable_params:,})")
47 print(f"~Param+optimizer VRAM (fp32): {human_bytes(param_mem_bytes)}")
48
49
50 # Your code to verify data ranges and integrity
51 # Hint: Create functions to check data ranges in training and validation data
52
53 # Enter your code here:
54 # ---- Data integrity checks ----
55 import torch.nn.functional as F
56 import torch.optim as optim # Added import for optimizers
57 from torch.utils.data import DataLoader
58 import numpy as np
59 import torch # Added import for torch
60
61 @torch.no_grad()
62 def inspect_loader(loader, name="loader", max_batches=3):
63     if loader is None:
64         print(f"{name}: not available")
65         return
66
67     mins, maxs, means, stds = [], [], [], []
68     labels = []
69     n_samples = 0
70
71     for i, (xb, yb) in enumerate(loader):
72         xb = xb.to(device)
73         mins.append(float(xb.min()))
74         maxs.append(float(xb.max()))
75         means.append(float(xb.mean()))
76         stds.append(float(xb.std()))
77         labels += yb.tolist()
78         n_samples += xb.size(0)
79         if i + 1 >= max_batches:
80             break
81
82     import numpy as np
83     if not mins:
84         print(f"{name}: (no batches)")
85         return
86
87     print(f"\n[{name}] ~{n_samples} samples inspected")
88     print(f" value min/max: {min(mins):.3f} / {max(maxs):.3f}")
89     print(f" mean ± std: {np.mean(means):.3f} ± {np.mean(stds):.3f}")
90     uniq = sorted(set(labels))
91     print(f" labels seen: {uniq[:20]}{'...' if len(uniq)>20 else ''}")

```

```

92
93 # Run checks (if you created these loaders earlier)
94 inspect_loader(globals().get("train_loader"), "train_loader")
95 inspect_loader(globals().get("val_loader"), "val_loader")
96
97
98 # Set up the optimizer with parameters tuned for diffusion models
99 # Note: Lower learning rates tend to work better for diffusion models
100 initial_lr = 0.001 # Starting learning rate
101 weight_decay = 1e-5 # L2 regularization to prevent overfitting
102
103 optimizer = optim.Adam( # Changed Adam to optim.Adam
104     model.parameters(),
105     lr=initial_lr,
106     weight_decay=weight_decay
107 )
108
109 # Learning rate scheduler to reduce LR when validation loss plateaus
110 # This helps fine-tune the model toward the end of training
111 # Compatible with all PyTorch versions
112 scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(
113     optimizer,
114     mode="min",           # Reduce LR when monitored value stops decreasing
115     factor=0.5,          # Multiply LR by this factor
116     patience=5,          # Number of epochs with no improvement
117     min_lr=1e-6          # Lower bound on the learning rate
118 )
119
120 # STUDENT EXPERIMENT:
121 # Try different channel configurations and see how they affect:
122 # 1. Model size (parameter count)
123 # 2. Training time
124 # 3. Generated image quality
125 #
126 # Suggestions:
127 # - Smaller: down_chs=(16, 32, 64)
128 # - Larger: down_chs=(64, 128, 256, 512)
129 def step_scheduler(val_loss):
130     old_lr = optimizer.param_groups[0]["lr"]
131     scheduler.step(val_loss)
132     new_lr = optimizer.param_groups[0]["lr"]
133     if new_lr != old_lr:
134         print(f"🌟 Learning rate reduced: {old_lr:.6f} → {new_lr:.6f}")

```

```

Created DownBlock: in_chs=32, out_chs=64, spatial_reduction=2x
Created DownBlock: in_chs=64, out_chs=128, spatial_reduction=2x
Created UpBlock: in_chs=128, out_chs=64, spatial_increase=2x
Created UpBlock: in_chs=64, out_chs=32, spatial_increase=2x
Created UNet with 3 scale levels
Channel dimensions: (32, 64, 128)

```

```

=====
MODEL ARCHITECTURE SUMMARY
=====
Input resolution: 28x28
Input channels: 1
Time steps: 100
Condition classes: 10
GPU acceleration: Yes

```

```

Parameters: 1,087,901 (trainable: 1,087,901)
~Param+optimizer VRAM (fp32): 16.60 MB

```

```

[train_loader] ~192 samples inspected
value min/max: -1.000 / 1.000
mean ± std: -0.734 ± 0.619
labels seen: [0, 1, 2, 3, 4, 5, 6, 7, 8, 9]
val_loader: not available

```

```

1 # Define helper functions needed for training and evaluation
2 def validate_model_parameters(model):
3     """
4         Counts model parameters and estimates memory usage.
5     """
6     total_params = sum(p.numel() for p in model.parameters())
7     trainable_params = sum(p.numel() for p in model.parameters() if p.requires_grad)
8
9     print(f"Total parameters: {total_params:,}")

```

```

10    print(f"Trainable parameters: {trainable_params:,}")
11
12    # Estimate memory requirements (very approximate)
13    param_memory = total_params * 4 / (1024 ** 2) # MB for params (float32)
14    grad_memory = trainable_params * 4 / (1024 ** 2) # MB for gradients
15    buffer_memory = param_memory * 2 # Optimizer state, forward activations, etc.
16
17    print(f"Estimated GPU memory usage: {param_memory + grad_memory + buffer_memory:.1f} MB")
18
19 # Define helper functions for verifying data ranges
20 def verify_data_range(dataloader, name="Dataset"):
21     """
22     Verifies the range and integrity of the data.
23     """
24     batch = next(iter(dataloader))[0]
25     print(f"\n{name} range check:")
26     print(f"Shape: {batch.shape}")
27     print(f"Data type: {batch.dtype}")
28     print(f"Min value: {batch.min().item():.2f}")
29     print(f"Max value: {batch.max().item():.2f}")
30     print(f"Contains NaN: {torch.isnan(batch).any().item()}")
31     print(f"Contains Inf: {torch.isinf(batch).any().item()}")
32
33 # Define helper functions for generating samples during training
34 from torchvision.utils import make_grid # Added import for make_grid
35 import matplotlib.pyplot as plt # Added import for plotting
36
37 def generate_samples(model, n_samples=10):
38     """
39     Generates sample images using the model for visualization during training.
40     """
41     model.eval()
42     with torch.no_grad():
43         # Generate digits 0-9 for visualization
44         samples = []
45         for digit in range(min(n_samples, 10)):
46             # Start with random noise
47             x = torch.randn(1, IMG_CH, IMG_SIZE, IMG_SIZE).to(device)
48
49             # Set up conditioning for the digit
50             c = torch.tensor([digit]).to(device)
51             c_one_hot = F.one_hot(c, N_CLASSES).float().to(device)
52             c_mask = torch.ones_like(c.unsqueeze(-1)).to(device)
53
54             # Remove noise step by step
55             for t in range(n_steps-1, -1, -1):
56                 t_batch = torch.full((1,), t).to(device)
57                 x = remove_noise(x, t_batch, model, c_one_hot, c_mask)
58
59             samples.append(x)
60
61         # Combine samples and display
62         samples = torch.cat(samples, dim=0)
63         grid = make_grid(samples, nrow=min(n_samples, 5), normalize=True)
64
65         plt.figure(figsize=(10, 4))
66
67         # Display based on channel configuration
68         if IMG_CH == 1:
69             plt.imshow(grid[0].cpu(), cmap='gray')
70         else:
71             plt.imshow(grid.permute(1, 2, 0).cpu())
72
73         plt.axis('off')
74         plt.title('Generated Samples')
75         plt.show()
76
77 # Define helper functions for safely saving models
78 def safe_save_model(model, path, optimizer=None, epoch=None, best_loss=None):
79     """
80     Safely saves model with error handling and backup.
81     """
82     try:
83         # Create a dictionary with all the elements to save
84         save_dict = {
85             'model_state_dict': model.state_dict(),
86         }

```

```

87     # Add optional elements if provided
88     if optimizer is not None:
89         save_dict['optimizer_state_dict'] = optimizer.state_dict()
90     if epoch is not None:
91         save_dict['epoch'] = epoch
92     if best_loss is not None:
93         save_dict['best_loss'] = best_loss
94
95     # Create a backup of previous checkpoint if it exists
96     if os.path.exists(path):
97         backup_path = path + '.backup'
98         try:
99             os.replace(path, backup_path)
100            print(f"Created backup at {backup_path}")
101        except Exception as e:
102            print(f"Warning: Could not create backup - {e}")
103
104    # Save the new checkpoint
105    torch.save(save_dict, path)
106    print(f"Model successfully saved to {path}")
107
108
109 except Exception as e:
110     print(f"Error saving model: {e}")
111     print("Attempting emergency save...")
112
113 try:
114     emergency_path = path + '.emergency'
115     torch.save(model.state_dict(), emergency_path)
116     print(f"Emergency save successful: {emergency_path}")
117 except:
118     print("Emergency save failed. Could not save model.")

```

```

1 # Implementation of the training step function
2 def train_step(x, c):
3     """
4     Performs a single training step for the diffusion model.
5     Returns a scalar tensor loss.
6     """
7     model.train()
8     x = x.to(device, non_blocking=True)
9     c = c.to(device, non_blocking=True)
10
11    # 1) Class conditioning (one-hot) + mask
12    c_one_hot = F.one_hot(c.long(), N_CLASSES).float().to(device)
13    c_mask = torch.ones(c.size(0), 1, device=device) # [B, 1]
14
15    # 2) Random timesteps for each sample
16    t = torch.randint(0, n_steps, (x.size(0),), device=device, dtype=torch.long)
17
18    # 3) Forward diffusion: add noise
19    x_t, true_noise = add_noise(x, t) # true_noise is the target
20
21    # 4) Predict noise
22    pred_noise = model(x_t, t, c_one_hot, c_mask)
23
24    # 5) Loss (MSE between predicted and actual noise)
25    loss = F.mse_loss(pred_noise, true_noise)
26
27    # 6) Optimize
28    optimizer.zero_grad(set_to_none=True)
29    loss.backward()
30    torch.nn.utils.clip_grad_norm_(model.parameters(), 1.0)
31    optimizer.step()
32
33    return loss.detach()
34

```

```

1 # Implementation of the main training loop
2 # Training configuration
3 early_stopping_patience = 10 # Number of epochs without improvement before stopping
4 gradient_clip_value = 1.0      # Maximum gradient norm for stability
5 display_frequency = 100        # How often to show progress (in steps)
6 generate_frequency = 500       # How often to generate samples (in steps)
7
8 # Progress tracking variables

```

```

9 best_loss = float('inf')
10 train_losses = []
11 val_losses = []
12 no_improve_epochs = 0
13
14 # Training loop
15 print("\n" + "="*50)
16 print("STARTING TRAINING")
17 print("=*50")
18 model.train()
19
20 # Wrap the training loop in a try-except block for better error handling:
21 # Your code for the training loop
22 # Hint: Use a try-except block for better error handling
23 # Process each epoch and each batch, with validation after each epoch
24
25 # Enter your code here:
26 import traceback
27 import matplotlib.pyplot as plt # Added import for plotting
28
29 # ---- Validation step (no grad, no optimizer) ----
30 @torch.no_grad()
31 def val_step(x, c):
32     model.eval()
33     x = x.to(device, non_blocking=True)
34     c = c.to(device, non_blocking=True)
35
36     c_one_hot = F.one_hot(c.long(), N_CLASSES).float().to(device)
37     c_mask = torch.ones(c.size(0), 1, device=device)
38
39     t = torch.randint(0, n_steps, (x.size(0),), device=device, dtype=torch.long)
40     x_t, true_noise = add_noise(x, t)
41     pred_noise = model(x_t, t, c_one_hot, c_mask)
42     return F.mse_loss(pred_noise, true_noise)
43
44 # ---- Main training loop ----
45 early_stopping_patience = 10
46 gradient_clip_value = 1.0      # (already handled inside train_step)
47 display_frequency = 100
48 generate_frequency = 500
49
50 best_loss = float('inf')
51 train_losses, val_losses = [], []
52 no_improve_epochs = 0
53
54 print("\n" + "="*50)
55 print("STARTING TRAINING")
56 print("=*50")
57
58 try:
59     for epoch in range(EPOCHS):
60         print(f"\nEpoch {epoch+1}/{EPOCHS}")
61         print("-" * 20)
62
63         # Training phase
64         model.train()
65         epoch_losses = []
66
67         for step, (images, labels) in enumerate(train_loader): # <- use train_loader
68             # Single training step (this already zero_grads, backward, clip, and step)
69             loss = train_step(images, labels)
70             epoch_losses.append(loss.item())
71
72             # Progress
73             if (step % display_frequency) == 0:
74                 print(f" Step {step}/{len(train_loader)} | Loss: {loss.item():.4f}")
75
76             # Occasional sampling
77             if step > 0 and (step % generate_frequency) == 0:
78                 print(" Generating samples...")
79                 generate_samples(model, n_samples=5)
80
81         # End of epoch: train summary
82         avg_train_loss = sum(epoch_losses) / max(1, len(epoch_losses))
83         train_losses.append(avg_train_loss)
84         print(f"\nTraining - Epoch {epoch+1} average loss: {avg_train_loss:.4f}")
85
86         # Validation phase

```

```

1 # validation phase
2 model.eval()
3 val_epoch_losses = []
4 print("Running validation...")
5 with torch.no_grad():
6     for val_images, val_labels in test_loader: # Changed val_loader to test_loader
7         vloss = val_step(val_images, val_labels)
8         val_epoch_losses.append(vloss.item())
9
10 avg_val_loss = sum(val_epoch_losses) / max(1, len(val_epoch_losses))
11 val_losses.append(avg_val_loss)
12 print(f"Validation - Epoch {epoch+1} average loss: {avg_val_loss:.4f}")
13
14 # LR scheduling
15 # If you defined step_scheduler wrapper earlier, use that; else use scheduler.step directly:
16 # step_scheduler(avg_val_loss)
17 scheduler.step(avg_val_loss)
18 current_lr = optimizer.param_groups[0]['lr']
19 print(f"Learning rate: {current_lr:.6f}")
20
21 # Epoch-end sampling
22 if (epoch % 2 == 0) or (epoch == EPOCHS - 1):
23     print("\nGenerating samples for visual progress check...")
24     generate_samples(model, n_samples=10)
25
26 # Save best
27 if avg_val_loss < best_loss:
28     best_loss = avg_val_loss
29     safe_save_model(model, 'best_diffusion_model.pt', optimizer, epoch, best_loss)
30     print(f"✓ New best model saved! (Val Loss: {best_loss:.4f})")
31     no_improve_epochs = 0
32 else:
33     no_improve_epochs += 1
34     print(f"No improvement for {no_improve_epochs}/{early_stopping_patience} epochs")
35
36 # Early stopping
37 if no_improve_epochs >= early_stopping_patience:
38     print("\nEarly stopping triggered! No improvement in validation loss.")
39     break
40
41 # Plot losses periodically
42 if (epoch % 5 == 0) or (epoch == EPOCHS - 1):
43     plt.figure(figsize=(10, 5))
44     plt.plot(train_losses, label='Training Loss')
45     plt.plot(val_losses, label='Validation Loss')
46     plt.xlabel('Epoch')
47     plt.ylabel('Loss')
48     plt.title('Training and Validation Loss')
49     plt.legend()
50     plt.grid(True)
51     plt.show()
52
53 except KeyboardInterrupt:
54     print("\n■ Training interrupted by user (KeyboardInterrupt).")
55 except Exception as e:
56     print("\n✖ Exception during training:")
57     print(e)
58     traceback.print_exc()
59 finally:
60     print("\n" + "="*50)
61     print("TRAINING COMPLETE")
62     print("=*50")
63     print(f"Best validation loss: {best_loss:.4f}")
64
65     print("Generating final samples...")
66     generate_samples(model, n_samples=10)
67
68 # Final loss curves
69 plt.figure(figsize=(12, 5))
70 plt.plot(train_losses, label='Training Loss')
71 plt.plot(val_losses, label='Validation Loss')
72 plt.xlabel('Epoch')
73 plt.ylabel('Loss')
74 plt.title('Training and Validation Loss')
75 plt.legend()
76 plt.grid(True)
77 plt.show()
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163

```

```
164     torch.cuda.empty_cache()
```

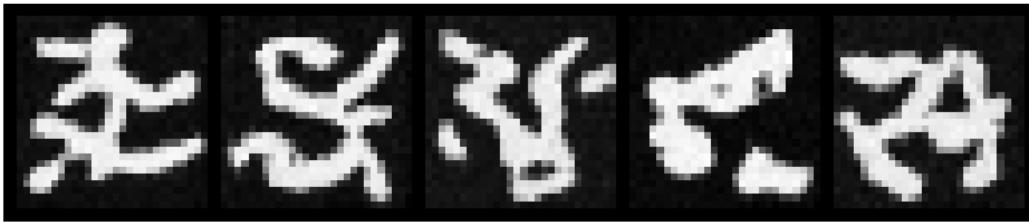


```
=====
STARTING TRAINING
=====

=====
STARTING TRAINING
=====

Epoch 1/30
-----
Step 0/938 | Loss: 0.0599
Step 100/938 | Loss: 0.0669
Step 200/938 | Loss: 0.0596
Step 300/938 | Loss: 0.0696
Step 400/938 | Loss: 0.0685
Step 500/938 | Loss: 0.0611
Generating samples...
```

Generated Samples



```
Step 600/938 | Loss: 0.0627
Step 700/938 | Loss: 0.0614
Step 800/938 | Loss: 0.0598
Step 900/938 | Loss: 0.0723
```

Training - Epoch 1 average loss: 0.0642

Running validation...

Validation - Epoch 1 average loss: 0.0642

Learning rate: 0.001000

Generating samples for visual progress check...

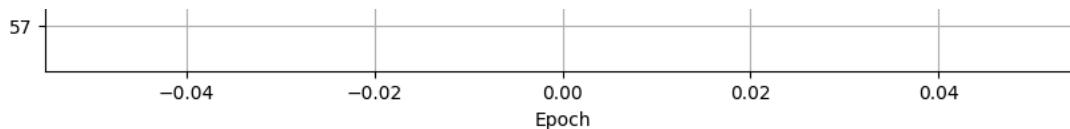
Generated Samples



Model successfully saved to best_diffusion_model.pt

✓ New best model saved! (Val Loss: 0.0642)





Epoch 2/30

```
Step 0/938 | Loss: 0.0611
Step 100/938 | Loss: 0.0593
Step 200/938 | Loss: 0.0684
Step 300/938 | Loss: 0.0714
Step 400/938 | Loss: 0.0672
Step 500/938 | Loss: 0.0659
Generating samples...
```

Generated Samples



```
Step 600/938 | Loss: 0.0650
Step 700/938 | Loss: 0.0658
Step 800/938 | Loss: 0.0589
Step 900/938 | Loss: 0.0544
```

Training - Epoch 2 average loss: 0.0627

Running validation...

Validation - Epoch 2 average loss: 0.0619

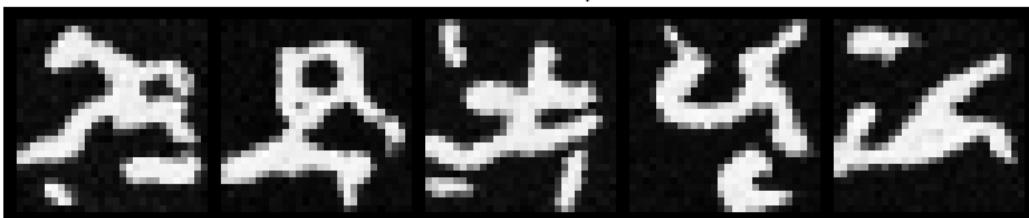
Learning rate: 0.001000

```
Created backup at best_diffusion_model.pt.backup
Model successfully saved to best_diffusion_model.pt
✓ New best model saved! (Val Loss: 0.0619)
```

Epoch 3/30

```
Step 0/938 | Loss: 0.0565
Step 100/938 | Loss: 0.0674
Step 200/938 | Loss: 0.0612
Step 300/938 | Loss: 0.0697
Step 400/938 | Loss: 0.0553
Step 500/938 | Loss: 0.0533
Generating samples...
```

Generated Samples



```
Step 600/938 | Loss: 0.0606
Step 700/938 | Loss: 0.0544
Step 800/938 | Loss: 0.0562
Step 900/938 | Loss: 0.0657
```

Training - Epoch 3 average loss: 0.0618

Running validation...

Validation - Epoch 3 average loss: 0.0621

Learning rate: 0.001000

Generating samples for visual progress check...

Generated Samples





No improvement for 1/10 epochs

Epoch 4/30

```
Step 0/938 | Loss: 0.0530
Step 100/938 | Loss: 0.0598
Step 200/938 | Loss: 0.0754
Step 300/938 | Loss: 0.0549
Step 400/938 | Loss: 0.0534
Step 500/938 | Loss: 0.0560
Generating samples...
```

Generated Samples



```
Step 600/938 | Loss: 0.0717
Step 700/938 | Loss: 0.0568
Step 800/938 | Loss: 0.0667
Step 900/938 | Loss: 0.0571
```

Training - Epoch 4 average loss: 0.0606

Running validation...

Validation - Epoch 4 average loss: 0.0614

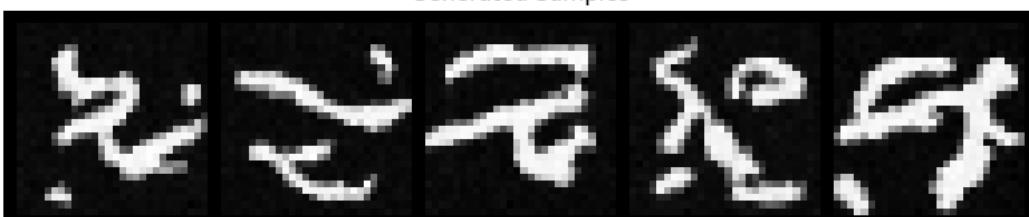
Learning rate: 0.001000

```
Created backup at best_diffusion_model.pt.backup
Model successfully saved to best_diffusion_model.pt
✓ New best model saved! (Val Loss: 0.0614)
```

Epoch 5/30

```
Step 0/938 | Loss: 0.0646
Step 100/938 | Loss: 0.0594
Step 200/938 | Loss: 0.0591
Step 300/938 | Loss: 0.0745
Step 400/938 | Loss: 0.0526
Step 500/938 | Loss: 0.0646
Generating samples...
```

Generated Samples



```
Step 600/938 | Loss: 0.0588
Step 700/938 | Loss: 0.0560
Step 800/938 | Loss: 0.0671
Step 900/938 | Loss: 0.0521
```

Training - Epoch 5 average loss: 0.0601

Running validation...

Validation - Epoch 5 average loss: 0.0607

Learning rate: 0.001000

Generating samples for visual progress check...

Generated Samples





Created backup at best_diffusion_model.pt.backup
 Model successfully saved to best_diffusion_model.pt
 ✓ New best model saved! (Val Loss: 0.0607)

Epoch 6/30

```
Step 0/938 | Loss: 0.0705
Step 100/938 | Loss: 0.0687
Step 200/938 | Loss: 0.0645
Step 300/938 | Loss: 0.0519
Step 400/938 | Loss: 0.0547
Step 500/938 | Loss: 0.0601
Generating samples...
```

Generated Samples



```
Step 600/938 | Loss: 0.0672
Step 700/938 | Loss: 0.0584
Step 800/938 | Loss: 0.0628
Step 900/938 | Loss: 0.0558
```

Training - Epoch 6 average loss: 0.0595

Running validation...

Validation - Epoch 6 average loss: 0.0587

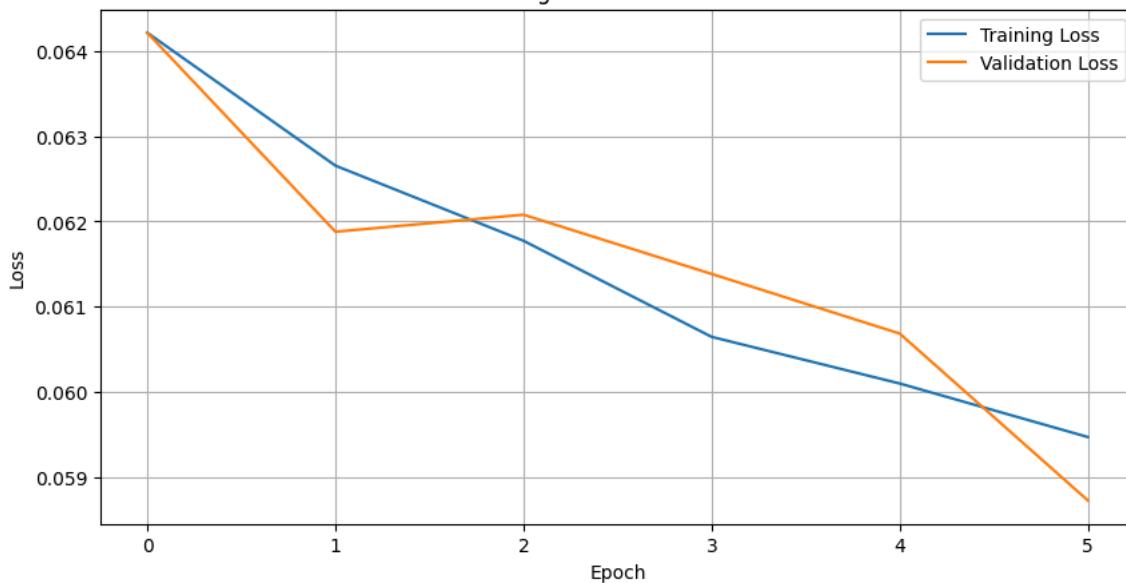
Learning rate: 0.001000

Created backup at best_diffusion_model.pt.backup

Model successfully saved to best_diffusion_model.pt

✓ New best model saved! (Val Loss: 0.0587)

Training and Validation Loss



Epoch 7/30

```
Step 0/938 | Loss: 0.0538
Step 100/938 | Loss: 0.0489
Step 200/938 | Loss: 0.0549
Step 300/938 | Loss: 0.0553
Step 400/938 | Loss: 0.0623
Step 500/938 | Loss: 0.0623
Generating samples...
```

Generated Samples



Step 600/938 | Loss: 0.0627
 Step 700/938 | Loss: 0.0550
 Step 800/938 | Loss: 0.0612
 Step 900/938 | Loss: 0.0580

Training - Epoch 7 average loss: 0.0591

Running validation...

Validation - Epoch 7 average loss: 0.0587

Learning rate: 0.001000

Generating samples for visual progress check...

Generated Samples



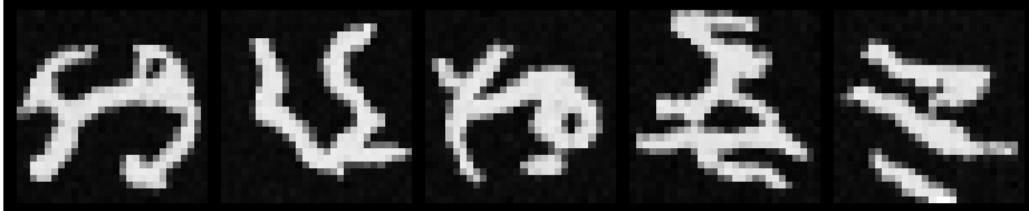
Created backup at best_diffusion_model.pt.backup
 Model successfully saved to best_diffusion_model.pt
 ✓ New best model saved! (Val Loss: 0.0587)

Epoch 8/30

Step 0/938 | Loss: 0.0512
 Step 100/938 | Loss: 0.0553
 Step 200/938 | Loss: 0.0610
 Step 300/938 | Loss: 0.0547
 Step 400/938 | Loss: 0.0525
 Step 500/938 | Loss: 0.0592

Generating samples...

Generated Samples



Step 600/938 | Loss: 0.0573
 Step 700/938 | Loss: 0.0741
 Step 800/938 | Loss: 0.0577
 Step 900/938 | Loss: 0.0496

Training - Epoch 8 average loss: 0.0589

Running validation...

Validation - Epoch 8 average loss: 0.0582

Learning rate: 0.001000

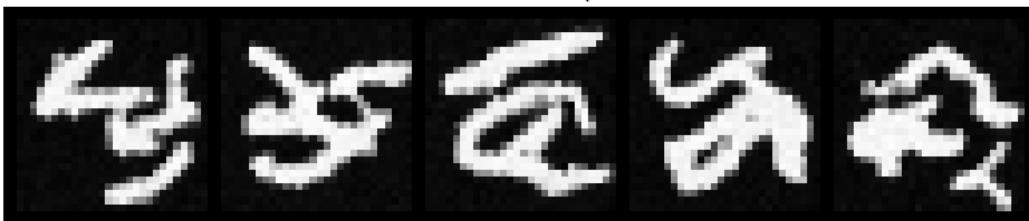
Created backup at best_diffusion_model.pt.backup
 Model successfully saved to best_diffusion_model.pt
 ✓ New best model saved! (Val Loss: 0.0582)

Epoch 9/30

Step 0/938 | Loss: 0.0552
 Step 100/938 | Loss: 0.0542

```
Step 200/938 | Loss: 0.0585
Step 300/938 | Loss: 0.0530
Step 400/938 | Loss: 0.0498
Step 500/938 | Loss: 0.0537
Generating samples...
```

Generated Samples



```
Step 600/938 | Loss: 0.0589
Step 700/938 | Loss: 0.0508
Step 800/938 | Loss: 0.0605
Step 900/938 | Loss: 0.0593
```

Training - Epoch 9 average loss: 0.0589

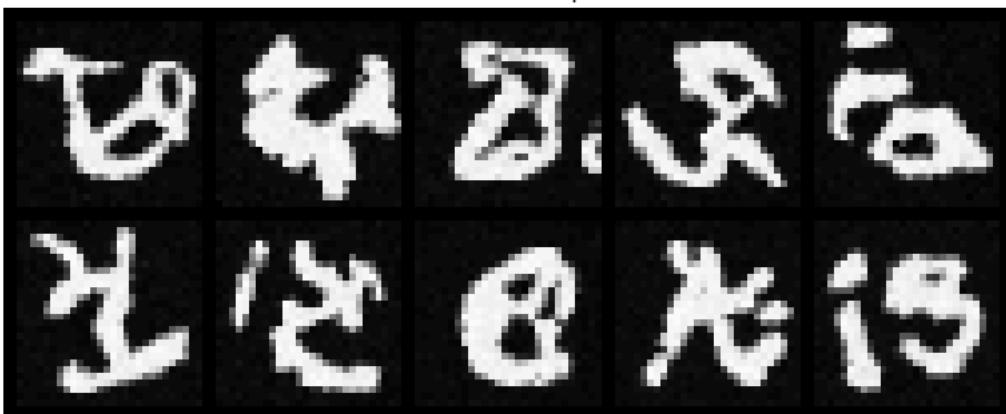
Running validation...

Validation - Epoch 9 average loss: 0.0597

Learning rate: 0.001000

Generating samples for visual progress check...

Generated Samples

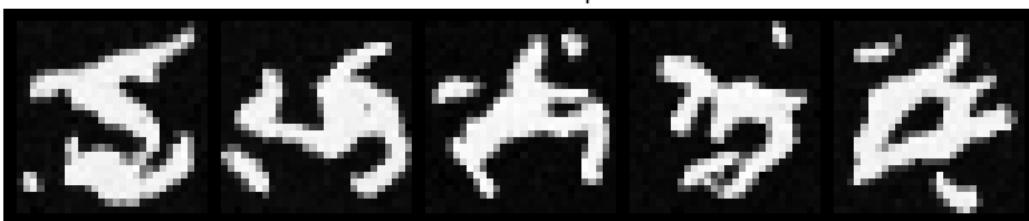


No improvement for 1/10 epochs

Epoch 10/30

```
Step 0/938 | Loss: 0.0552
Step 100/938 | Loss: 0.0540
Step 200/938 | Loss: 0.0626
Step 300/938 | Loss: 0.0516
Step 400/938 | Loss: 0.0673
Step 500/938 | Loss: 0.0548
Generating samples...
```

Generated Samples



```
Step 600/938 | Loss: 0.0577
Step 700/938 | Loss: 0.0611
Step 800/938 | Loss: 0.0616
Step 900/938 | Loss: 0.0613
```

Training - Epoch 10 average loss: 0.0583

Running validation...

Validation - Epoch 10 average loss: 0.0594

Learning rate: 0.001000

No improvement for 2/10 epochs

Epoch 11/30

```
Step 0/938 | Loss: 0.0512
```

```
Step 100/938 | Loss: 0.0526
Step 200/938 | Loss: 0.0556
Step 300/938 | Loss: 0.0545
Step 400/938 | Loss: 0.0536
Step 500/938 | Loss: 0.0590
Generating samples...
```

Generated Samples

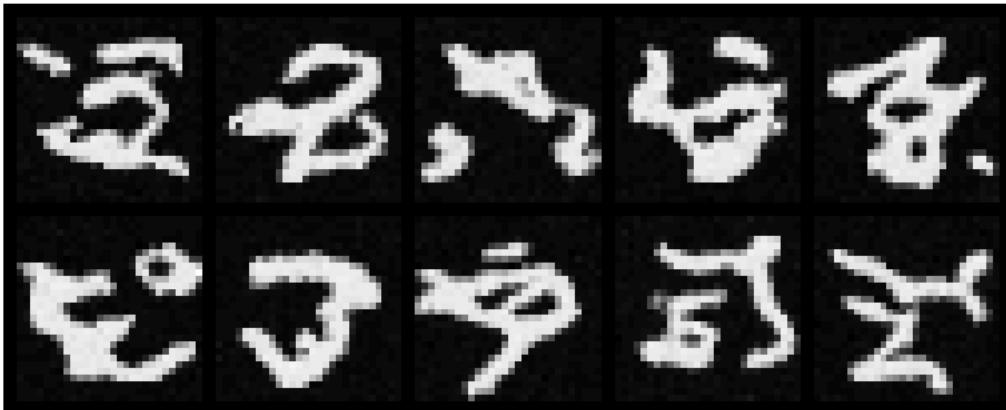


```
Step 600/938 | Loss: 0.0691
Step 700/938 | Loss: 0.0659
Step 800/938 | Loss: 0.0568
Step 900/938 | Loss: 0.0495
```

Training - Epoch 11 average loss: 0.0579
 Running validation...
 Validation - Epoch 11 average loss: 0.0576
 Learning rate: 0.001000

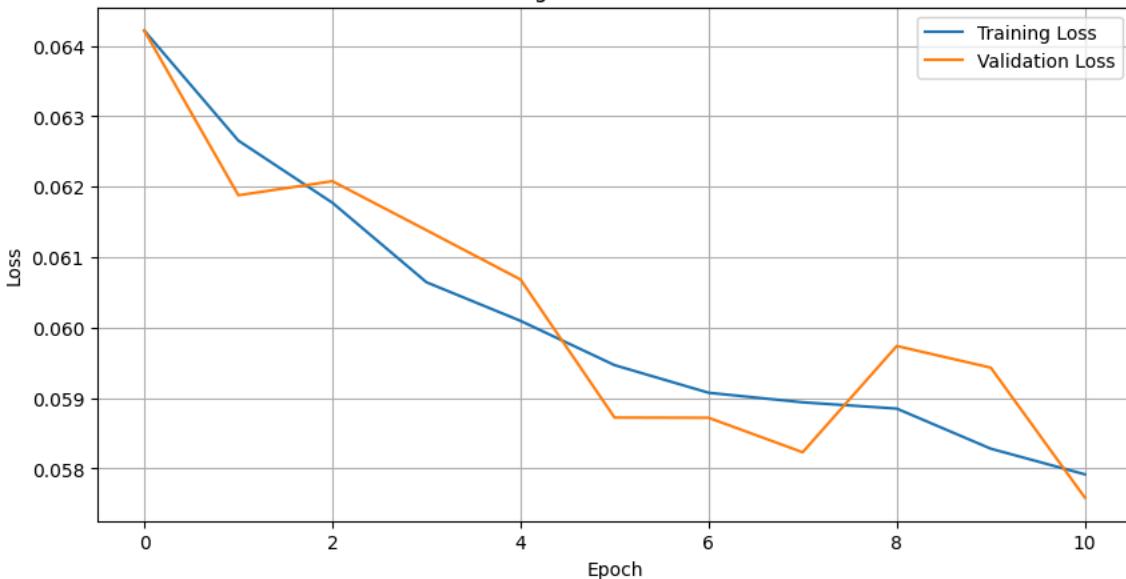
Generating samples for visual progress check...

Generated Samples



Created backup at best_diffusion_model.pt.backup
 Model successfully saved to best_diffusion_model.pt
 ✓ New best model saved! (Val Loss: 0.0576)

Training and Validation Loss



Epoch 12/30

```
Step 0/938 | Loss: 0.0636
Step 100/938 | Loss: 0.0568
Step 200/938 | Loss: 0.0634
```

```
Step 300/938 | Loss: 0.0553
Step 400/938 | Loss: 0.0579
Step 500/938 | Loss: 0.0564
Generating samples...
```

Generated Samples



```
Step 600/938 | Loss: 0.0623
Step 700/938 | Loss: 0.0573
Step 800/938 | Loss: 0.0531
Step 900/938 | Loss: 0.0625
```

Training - Epoch 12 average loss: 0.0579

Running validation...

Validation - Epoch 12 average loss: 0.0575

Learning rate: 0.001000

Created backup at best_diffusion_model.pt.backup

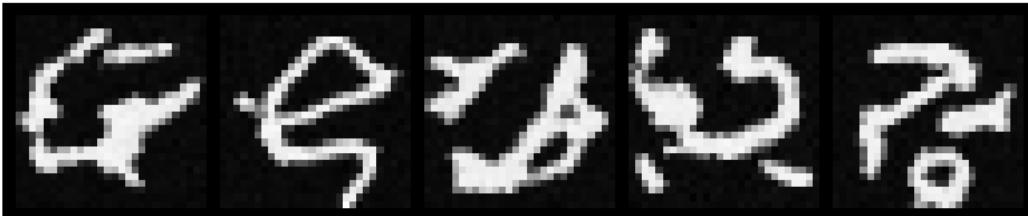
Model successfully saved to best_diffusion_model.pt

✓ New best model saved! (Val Loss: 0.0575)

Epoch 13/30

```
Step 0/938 | Loss: 0.0495
Step 100/938 | Loss: 0.0525
Step 200/938 | Loss: 0.0618
Step 300/938 | Loss: 0.0503
Step 400/938 | Loss: 0.0591
Step 500/938 | Loss: 0.0516
Generating samples...
```

Generated Samples



```
Step 600/938 | Loss: 0.0553
Step 700/938 | Loss: 0.0653
Step 800/938 | Loss: 0.0605
Step 900/938 | Loss: 0.0582
```

Training - Epoch 13 average loss: 0.0576

Running validation...

Validation - Epoch 13 average loss: 0.0574

Learning rate: 0.001000

Generating samples for visual progress check...

Generated Samples

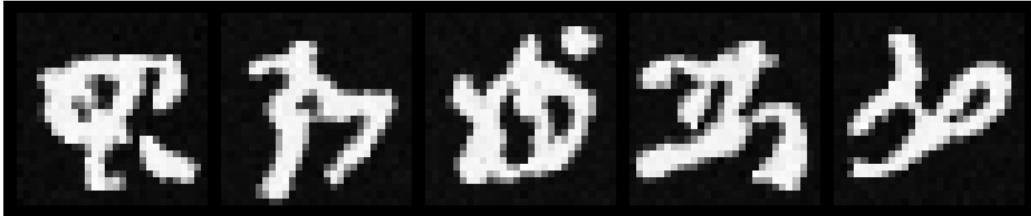


```
Created backup at best_diffusion_model.pt.backup
Model successfully saved to best_diffusion_model.pt
✓ New best model saved! (Val Loss: 0.0574)
```

Epoch 14/30

```
Step 0/938 | Loss: 0.0593
Step 100/938 | Loss: 0.0670
Step 200/938 | Loss: 0.0548
Step 300/938 | Loss: 0.0473
Step 400/938 | Loss: 0.0561
Step 500/938 | Loss: 0.0650
Generating samples...
```

Generated Samples



```
Step 600/938 | Loss: 0.0577
Step 700/938 | Loss: 0.0570
Step 800/938 | Loss: 0.0594
Step 900/938 | Loss: 0.0663
```

Training - Epoch 14 average loss: 0.0573

Running validation...

Validation - Epoch 14 average loss: 0.0570

Learning rate: 0.001000

Created backup at best_diffusion_model.pt.backup

Model successfully saved to best_diffusion_model.pt

✓ New best model saved! (Val Loss: 0.0570)

Epoch 15/30

```
Step 0/938 | Loss: 0.0534
Step 100/938 | Loss: 0.0665
Step 200/938 | Loss: 0.0606
Step 300/938 | Loss: 0.0614
Step 400/938 | Loss: 0.0475
Step 500/938 | Loss: 0.0569
Generating samples...
```

Generated Samples



```
Step 600/938 | Loss: 0.0561
Step 700/938 | Loss: 0.0630
Step 800/938 | Loss: 0.0503
Step 900/938 | Loss: 0.0543
```

Training - Epoch 15 average loss: 0.0574

Running validation...

Validation - Epoch 15 average loss: 0.0567

Learning rate: 0.001000

Generating samples for visual progress check...

Generated Samples

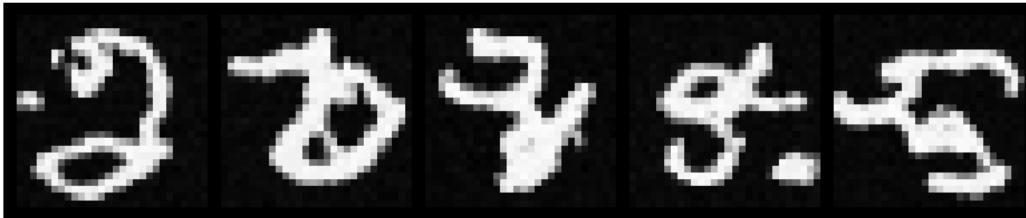


```
Created backup at best_diffusion_model.pt.backup
Model successfully saved to best_diffusion_model.pt
✓ New best model saved! (Val Loss: 0.0567)
```

Epoch 16/30

```
Step 0/938 | Loss: 0.0647
Step 100/938 | Loss: 0.0561
Step 200/938 | Loss: 0.0565
Step 300/938 | Loss: 0.0669
Step 400/938 | Loss: 0.0562
Step 500/938 | Loss: 0.0500
Generating samples...
```

Generated Samples



```
Step 600/938 | Loss: 0.0591
Step 700/938 | Loss: 0.0469
Step 800/938 | Loss: 0.0580
Step 900/938 | Loss: 0.0587
```

Training - Epoch 16 average loss: 0.0567

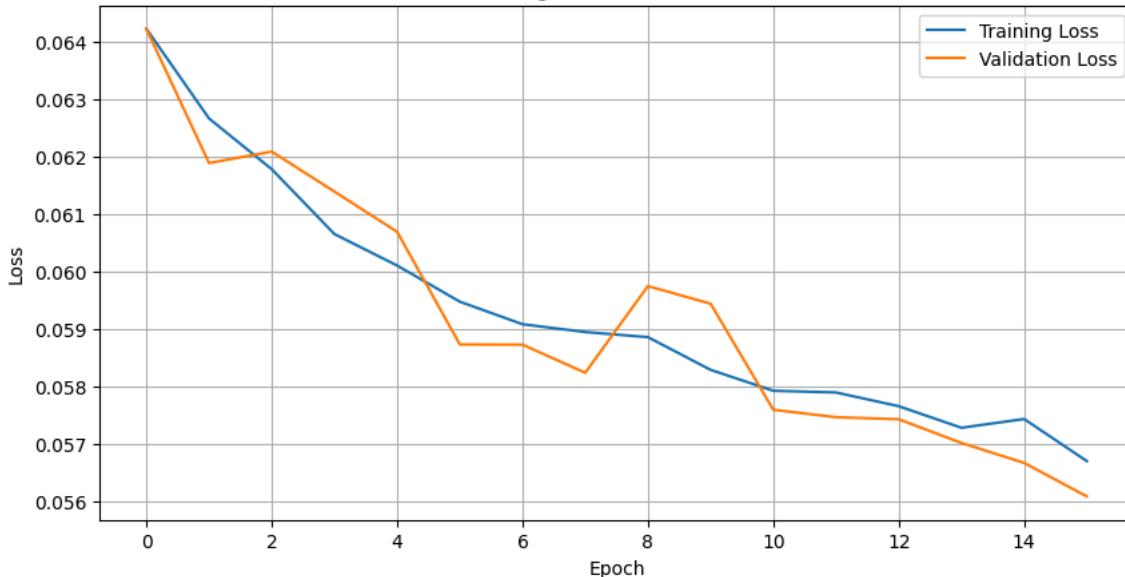
Running validation...

Validation - Epoch 16 average loss: 0.0561

Learning rate: 0.001000

```
Created backup at best_diffusion_model.pt.backup
Model successfully saved to best_diffusion_model.pt
✓ New best model saved! (Val Loss: 0.0561)
```

Training and Validation Loss



Epoch 17/30

```
Step 0/938 | Loss: 0.0572
Step 100/938 | Loss: 0.0552
Step 200/938 | Loss: 0.0614
Step 300/938 | Loss: 0.0618
Step 400/938 | Loss: 0.0480
Step 500/938 | Loss: 0.0574
Generating samples...
```

Generated Samples

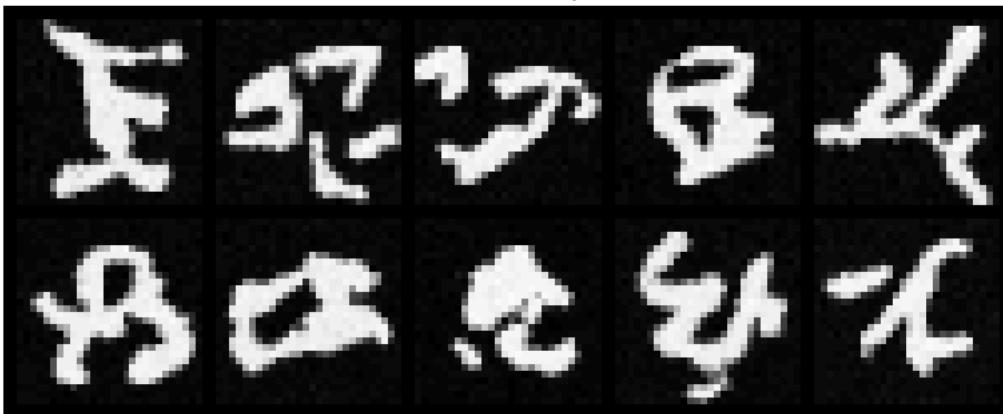


```
Step 600/938 | Loss: 0.0556
Step 700/938 | Loss: 0.0572
Step 800/938 | Loss: 0.0512
Step 900/938 | Loss: 0.0542
```

Training - Epoch 17 average loss: 0.0569
 Running validation...
 Validation - Epoch 17 average loss: 0.0561
 Learning rate: 0.001000

Generating samples for visual progress check...

Generated Samples



No improvement for 1/10 epochs

Epoch 18/30

```
Step 0/938 | Loss: 0.0635
Step 100/938 | Loss: 0.0544
Step 200/938 | Loss: 0.0574
Step 300/938 | Loss: 0.0565
Step 400/938 | Loss: 0.0536
Step 500/938 | Loss: 0.0528
Generating samples...
```

Generated Samples



```
Step 600/938 | Loss: 0.0523
Step 700/938 | Loss: 0.0540
Step 800/938 | Loss: 0.0591
Step 900/938 | Loss: 0.0598
```

Training - Epoch 18 average loss: 0.0569
 Running validation...
 Validation - Epoch 18 average loss: 0.0556
 Learning rate: 0.001000
 Created backup at best_diffusion_model.pt.backup
 Model successfully saved to best_diffusion_model.pt
 ✓ New best model saved! (Val Loss: 0.0556)

Epoch 19/30

```
Step 0/938 | Loss: 0.0636
Step 100/938 | Loss: 0.0574
Step 200/938 | Loss: 0.0602
Step 300/938 | Loss: 0.0615
Step 400/938 | Loss: 0.0606
Step 500/938 | Loss: 0.0538
Generating samples...
```

Generated Samples



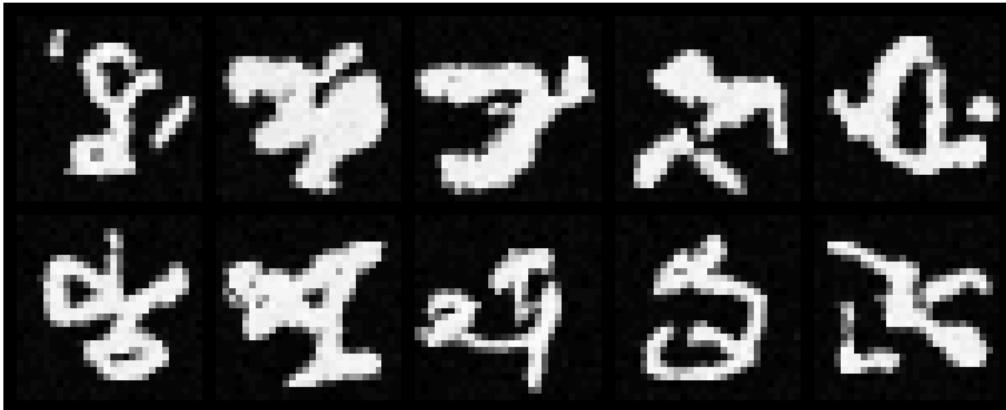


Step 600/938 | Loss: 0.0550
 Step 700/938 | Loss: 0.0523
 Step 800/938 | Loss: 0.0533
 Step 900/938 | Loss: 0.0572

Training - Epoch 19 average loss: 0.0568
 Running validation...
 Validation - Epoch 19 average loss: 0.0561
 Learning rate: 0.001000

Generating samples for visual progress check...

Generated Samples



No improvement for 1/10 epochs

Epoch 20/30

Step 0/938 | Loss: 0.0506
 Step 100/938 | Loss: 0.0621
 Step 200/938 | Loss: 0.0486
 Step 300/938 | Loss: 0.0485
 Step 400/938 | Loss: 0.0557
 Step 500/938 | Loss: 0.0467
 Generating samples...

Generated Samples



Step 600/938 | Loss: 0.0616
 Step 700/938 | Loss: 0.0552
 Step 800/938 | Loss: 0.0648
 Step 900/938 | Loss: 0.0544

Training - Epoch 20 average loss: 0.0565
 Running validation...
 Validation - Epoch 20 average loss: 0.0554
 Learning rate: 0.001000
 Created backup at best_diffusion_model.pt.backup
 Model successfully saved to best_diffusion_model.pt
 ✓ New best model saved! (Val Loss: 0.0554)

Epoch 21/30

Step 0/938 | Loss: 0.0558
 Step 100/938 | Loss: 0.0606
 Step 200/938 | Loss: 0.0575
 Step 300/938 | Loss: 0.0463
 Step 400/938 | Loss: 0.0482
 Step 500/938 | Loss: 0.0618
 Generating samples...

Generated Samples